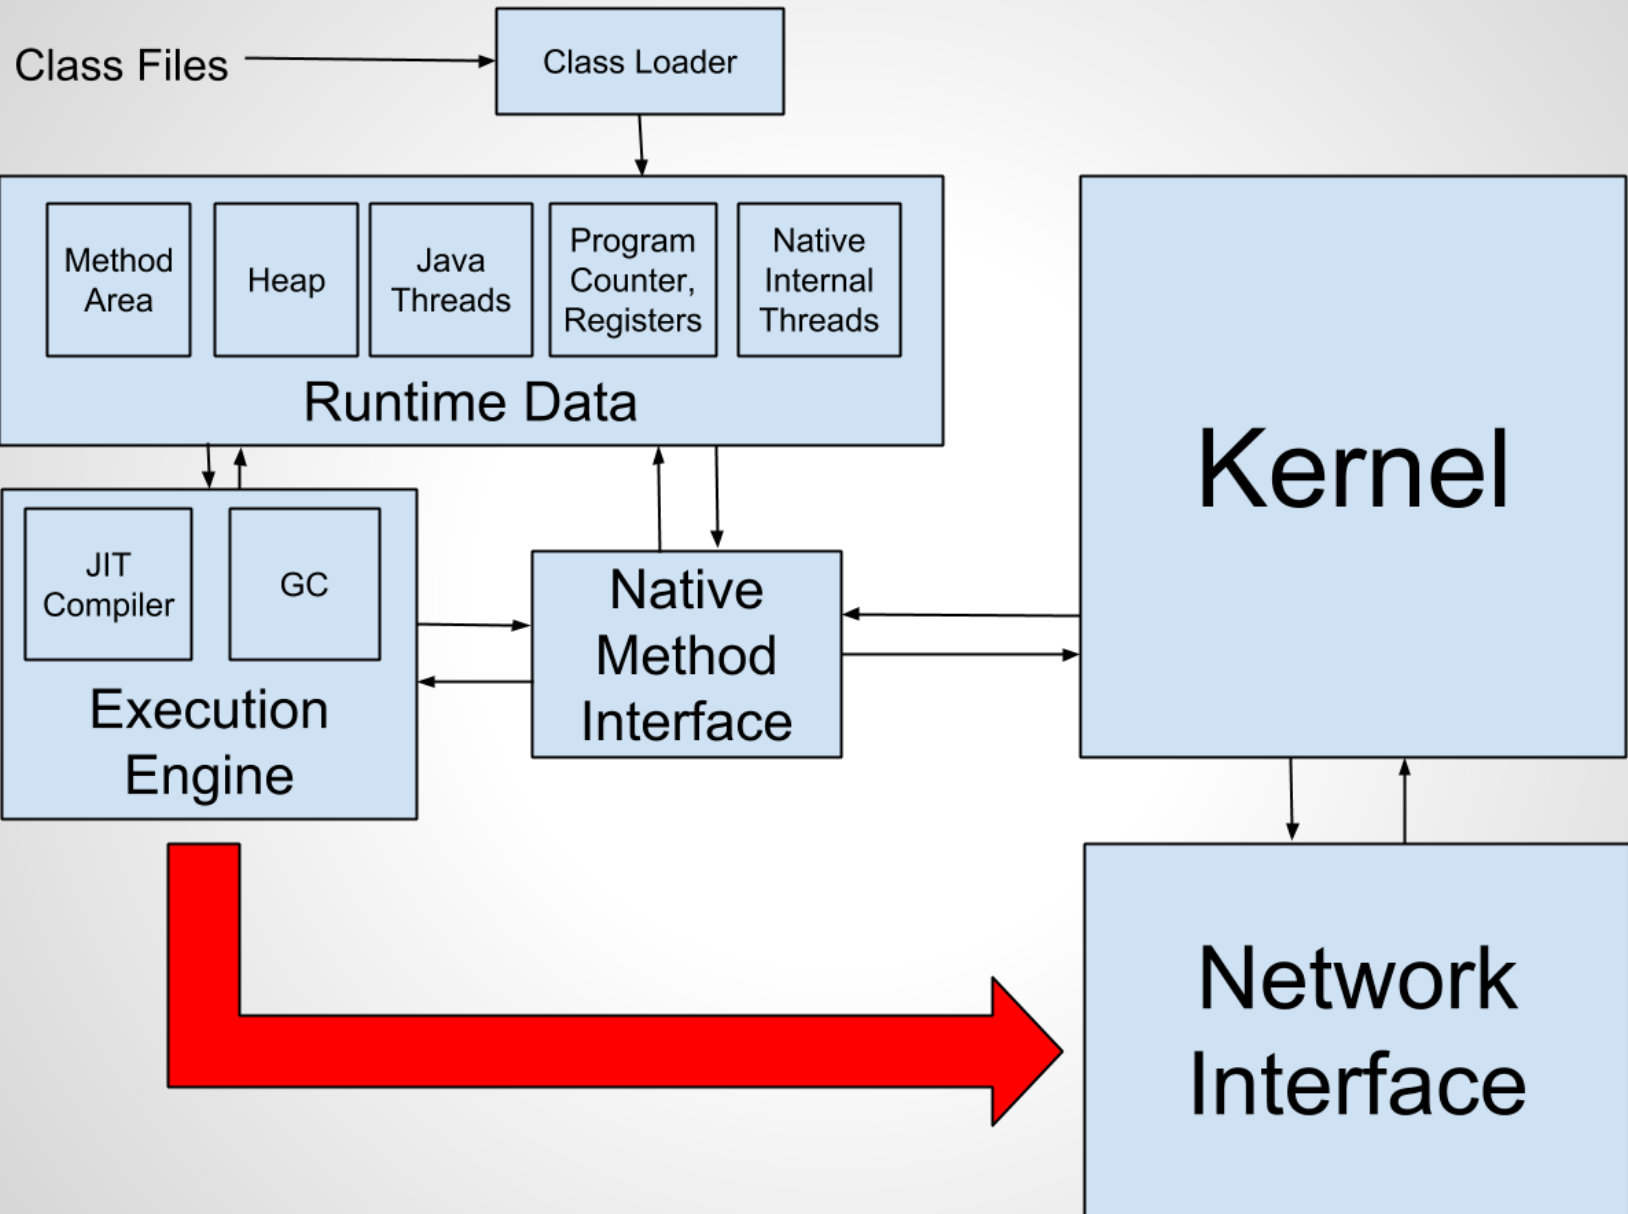


Packet Processing in Java

Ashley Hemingway

Whats the Problem?

- Middleboxes in native languages are inflexible to emerging technologies
- Poor Java networking capabilities even though high performance processing is possible



DPDK

- NIC driver in user space
- Idea of logical cores (lcores)
- Multiple queues per port
- No interrupts, just polling
- 2 application models

DPDK-Java

- Framework wrapping DPDK
- Easy creation of middlebox applications in Java
- Eliminates error checking
- Integration with other Java based applications

```

#include <stdio.h>
#include <string.h>
#include <stdint.h>
#include <errno.h>
#include <assert.h>
#include <sys/queue.h>

#include <rte_ip.h>
#include <rte_memory.h>
#include <rte_memzone.h>
#include <rte_launch.h>
#include <rte_tailq.h>
#include <rte_eal.h>
#include <rte_per_lcore.h>
#include <rte_lcore.h>
#include <rte_debug.h>
#include <rte_ethdev.h>
#include <rte_ring.h>
#include <rte_mempool.h>
#include <rte_mbuf.h>
#include <rte_timer.h>
#include <rte_cycles.h>

#include "main.h"

#define MAX_PKT_BURST 512

#define MAX_PKT_SIZE (2*1024 + sizeof
(struct rte_mbuf) +
RTE_PKTMBUF_HEADROOM)

#define NB_RX_QUEUE 2

#define NB_TX_QUEUE 2

#define NB_RX_DESC 256

#define NB_TX_DESC 256

struct rte_mbuf *rx_mbufs[MAX_PKT_BURST];

// For configuring an ethernet port
static struct rte_eth_conf eth_conf = {
    .rxmode = {
        .mq_mode = ETH_MQ_RX_RSS,
        .split_hdr_size = 0,
        .header_split = 0,
        .hw_ip_checksum = 1,
        .hw_vlan_filter = 0,
        .jumbo_frame = 0,
        .hw_strip_crc = 0,
    },
    .rx_adv_conf = {
        .rss_conf = {
            .rss_key = NULL,
            .rss_hf = ETH_RSS_IP,
        },
    },
    .txmode = {
        .mq_mode = ETH_MQ_TX_VMDQ_ONLY
    },
};

// Configures a TX ring of an Ethernet port
static struct rte_eth_txconf tx_conf = {
    .tx_thresh = {
        .pthresh = 36,
        .hthresh = 0,
        .wthresh = 0,
    },
    .tx_rs_thresh = 0,
    .tx_free_thresh = 0,
};

// Configures a RX ring of an Ethernet port
static struct rte_eth_rxconf rx_conf = {
    .rx_thresh = {
        .pthresh = 8,
        .hthresh = 8,
        .wthresh = 4,
    },
    .rx_free_thresh = 64,
    .rx_drop_en = 0,
};

struct mbuf_list {
    unsigned len;
    struct rte_mbuf *list[MAX_PKT_BURST];
} free_list, send_list;

// RTE mempool structure
static struct rte_mempool *rx_pool;
static void send_packet(struct rte_mbuf *);
static void free_packet(struct rte_mbuf *);
static void send_burst(void);
//static void free_burst(void);

static uint32_t blacklist[] = {0, 1};
int size = 2;

static void send_burst(void) {
    struct rte_mbuf **list = (struct rte_mbuf **)
send_list.list;
    int i;
    if (rte_lcore_id() == 1) {
        i = 0;
    } else {
        i = 1;
    }
    int ret = rte_eth_tx_burst(0, i, list,
MAX_PKT_BURST);
    if (unlikely(ret < MAX_PKT_BURST)) {
        do {
            free_packet(list[ret]);
        } while (++ret < MAX_PKT_BURST);
    }
}

static void send_packet(struct rte_mbuf *m) {
    unsigned len = send_list.len;
    send_list.list[len] = m;
    len++;

    /* enough pkts to be sent */
    if (unlikely(len == MAX_PKT_BURST)) {
        send_burst();
        len = 0;
    }
    send_list.len = len;
}

static void free_packet(struct rte_mbuf *m) {
    int main_loop(void *);

    int main_loop(__attribute__((unused)) void *arg) {
        // long pktcount = 0;
        int rcv_cnt, i, ifidx = 0;
        struct ipv4_hdr *iphdr;
        struct rte_mbuf *m;
        int id = rte_lcore_id();
        printf("Core %i starting\n", id);
        fflush(stdout);
        int b;
        if (id == 1) {
            b = 0;
        } else {
            b = 1;
        }
        while(1) {
            rcv_cnt = rte_eth_rx_burst(ifidx, b, rx_mbufs,
MAX_PKT_BURST);
            if (rcv_cnt < 0) {
                if (errno != EAGAIN && errno != EINTR) {
                    perror("rte_eth_rx_burst()");
                    assert(0);
                }
            }
            if ( rcv_cnt > 0) {
                pktcount += rcv_cnt;
                for (i = 0 ; i < rcv_cnt; i++) {
                    m = rx_mbufs[i];
                    iphdr = (struct ipv4_hdr *)rte_pktmbuf_adj
(m, (uint16_t)sizeof(struct ether_hdr));
                    //RTE_MBUF_ASSERT(iphdr != NULL);

                    uint32_t dest_addr = rte_be_to_cpu_32
(iphdr->dst_addr);

                    int drop = 0;

                    int a;
                    for (a = 0; a < size; a++) {
                        if (blacklist[a] == dest_addr) {
                            drop = 1;
                        }
                    }
                    if (drop == 0) {
                        send_packet(m);
                        continue;
                    }
                    free_packet(m);
                }
            }
            return 0;
        }
    }

    void timer_setup(void);
    void do_stats(struct rte_timer *, void *);

    //received bytes
    uint64_t pre_ibytes = 0;
    //received packets - successful
    uint64_t pre_ipackets = 0;

    void do_stats(__attribute__((unused)) struct
rte_mbuf *m) {
        void do_stats(struct rte_timer *, void *) {
            struct rte_timer timer;

            void timer_setup(void) {
                int lcore_id = rte_get_master_lcore();
                rte_timer_subsystem_init();
                rte_timer_init(&timer);
                int ret = rte_timer_reset(&timer,
rte_get_timer_hz(), PERIODICAL, lcore_id,
do_stats, NULL);
                if (ret != 0) {
                    printf("TIMER ERROR");
                    fflush(stdout);
                }
            }

            int main(int argc, char **argv) {
                int ret, ifidx = 0;
                //unsigned lcore_id;
                uint8_t count;

                struct rte_eth_link link;

                // initialise eal
                ret = rte_eal_init(argc, argv);
                if (ret < 0) {
                    rte_panic("Cannot init EAL\n");
                }

                count = rte_eth_dev_count();
                printf("# of eth ports = %d\n", count);
                //memset(&eth_conf, 0, sizeof eth_conf);
                ret = rte_eth_dev_configure(0, 2, 2, &eth_conf);
                if (ret < 0) {
                    rte_exit(EXIT_FAILURE, "Cannot configure
device: error=%d, port=%d\n", ret, 0);
                }
                printf("If %d rte_eth_dev_configure()
successful\n", ifidx);

                rx_pool = rte_mempool_create("rx_pool",
16*1024, MAX_PKT_SIZE, 0,
sizeof (struct
rte_pktmbuf_pool_private),
rte_pktmbuf_pool_init, NULL,
rte_pktmbuf_init, NULL, 1, 0);

                if (rx_pool == NULL) {
                    rte_exit(EXIT_FAILURE,
"rte_mempool_create(): error\n");
                }

                ret = rte_eth_rx_queue_setup(ifidx, 0, NB_RX_DESC, 1,
&rx_conf, rx_pool);
                if (ret < -1) {
                    rte_exit(EXIT_FAILURE,
"rte_eth_rx_queue_setup(): error=%d, port=%d\n", ret,
ifidx);
                }

                ret = rte_eth_tx_queue_setup(0, 1, NB_RX_DESC, 1,
&tx_conf);
                if (ret < 0) {
                    rte_exit(EXIT_FAILURE, "rte_eth_tx_queue_setup():
error=%d, port=%d\n", ret, ifidx);
                }
                printf("If %d rte_eth_tx_queue_setup()
successful\n", ifidx);

                ret = rte_eth_dev_start(ifidx);
                if (ret < 0) {
                    rte_exit(EXIT_FAILURE, "rte_eth_dev_start(): error=%
d, port=%d\n", ret, ifidx);
                }
                printf("If %d rte_eth_dev_start()
successful\n", ifidx);

                rte_eth_link_get(ifidx, &link);
                if (link.link_status == 0) {
                    rte_exit(EXIT_FAILURE, "DPDK interface is down: %
d\n", ifidx);
                }
                printf("If %d is UP and RUNNING\n", ifidx);

                rte_eth_promiscuous_enable(ifidx);

                struct lcore_config lc = lcore_config[0];
                printf("%i", lc.core_id);

                uint32_t a;
                for (a = 0; a < 32; a++) {
                    if (a == rte_get_master_lcore() || !rte_lcore_is_enabled
(a)) {
                        continue;
                    }
                    ret = rte_eal_remote_launch(main_loop, NULL, a);
                    if (ret != 0) {
                        printf("ERROR");
                    }
                }
                rte_delay_ms(1000);
                timer_setup();

                for(;;) {
                    rte_timer_manage();
                }
            }
        }
    }
}

```

```

import java.io.FileInputStream;
import java.io.FileNotFoundException;
import java.util.ArrayList;
import java.util.List;

public class Firewall {

    public static void main(String[] args) {

        ApplicationStarter as = new ApplicationStarter();

        try {
            as.readConfig(new FileInputStream("config.properties"));
        } catch (FileNotFoundException e) {
            e.printStackTrace();
        }

        as.sendDPDKInformation();

        List<CoreThread> threads = new ArrayList<CoreThread>();
        ReceivePoller rp1 = new ReceivePoller(0, 0);
        ReceivePoller rp2 = new ReceivePoller(0, 1);
        PacketSender ps1 = new PacketSender(0, 0);
        PacketSender ps2 = new PacketSender(0, 1);
        PacketFreer pf1 = new PacketFreer();
        PacketFreer pf2 = new PacketFreer();

        threads.add(new FirewallProcessor(ps1, pf1, rp1, "PROCESS 1", 3));
        threads.add(new FirewallProcessor(ps2, pf2, rp2, "PROCESS 2", 5));

        as.createAffinityThreads(threads);

        as.dpd_k_init_eal();
        as.dpd_k_create_mempool("mbufs", 8192*4, 32, 1);
        as.dpd_k_check_ports();
        as.dpd_k_configure_dev(0, 2, 2);
        as.dpd_k_configure_rx_queue(0, 0, 1);
        as.dpd_k_configure_tx_queue(0, 0, 1);
        as.dpd_k_configure_rx_queue(0, 1, 1);
        as.dpd_k_configure_tx_queue(0, 1, 1);
        as.dpd_k_dev_start(0);
        as.dpd_k_check_ports_link_status();

        as.dpd_k_enable_pro();
        as.start_native_stats(1);
        as.dpd_k_get_mac_info();

        as.startAll();
    }

import java.io.BufferedReader;
import java.io.File;
import java.io.FileNotFoundException;
import java.io.FileReader;
import java.io.IOException;
import java.util.HashSet;
import java.util.Set;

public class FirewallProcessor extends PacketProcessor {

    Set<Long> blacklist;

    PacketFreer pf_ind;
    ReceivePoller rp_ind;
    PacketSender ps_ind;

```

```

public FirewallProcessor(PacketSender ps, PacketFreer pf, ReceivePoller rp, int core) {
    super(ps, pf, rp, core);
    blacklist = new HashSet<Long>();
    readBlacklist();
    pf_ind = pf;
    rp_ind = rp;
    ps_ind = ps;
}

public FirewallProcessor(PacketSender ps, PacketFreer pf, ReceivePoller rp, String name, int core) {
    super(ps, pf, rp, name, core);
    blacklist = new HashSet<Long>();
    readBlacklist();
    pf_ind = pf;
    rp_ind = rp;
    ps_ind = ps;
}

private boolean inspect(Packet currentPacket) {
    int version = currentPacket.whichIP();

    if (version == 4) {
        Ipv4Packet cp = (Ipv4Packet)currentPacket;
        if (blacklist.contains(cp.getSrcAddr())) {
            pf_ind.freePacket(currentPacket);
        } else {
            ps_ind.sendPacket(currentPacket);
            return true;
        }
    } else if (version == 6) {
        pf_ind.freePacket(currentPacket);
    } else {
        pf_ind.freePacket(currentPacket);
    }
    return false;
}

private void readBlacklist() {
    File file = new File("blacklist.txt");
    FileReader fileReader;
    try {
        fileReader = new FileReader(file);
        BufferedReader bufferedReader = new BufferedReader(fileReader);
        String line;
        try {
            while ((line = bufferedReader.readLine()) != null) {
                blacklist.add(Utils.IpToInt(line));
            }
        } catch (IOException e1) {
            e1.printStackTrace();
        }
        try {
            fileReader.close();
        } catch (IOException e) {
            e.printStackTrace();
        }
    } catch (FileNotFoundException e) {
        e.printStackTrace();
    }
}

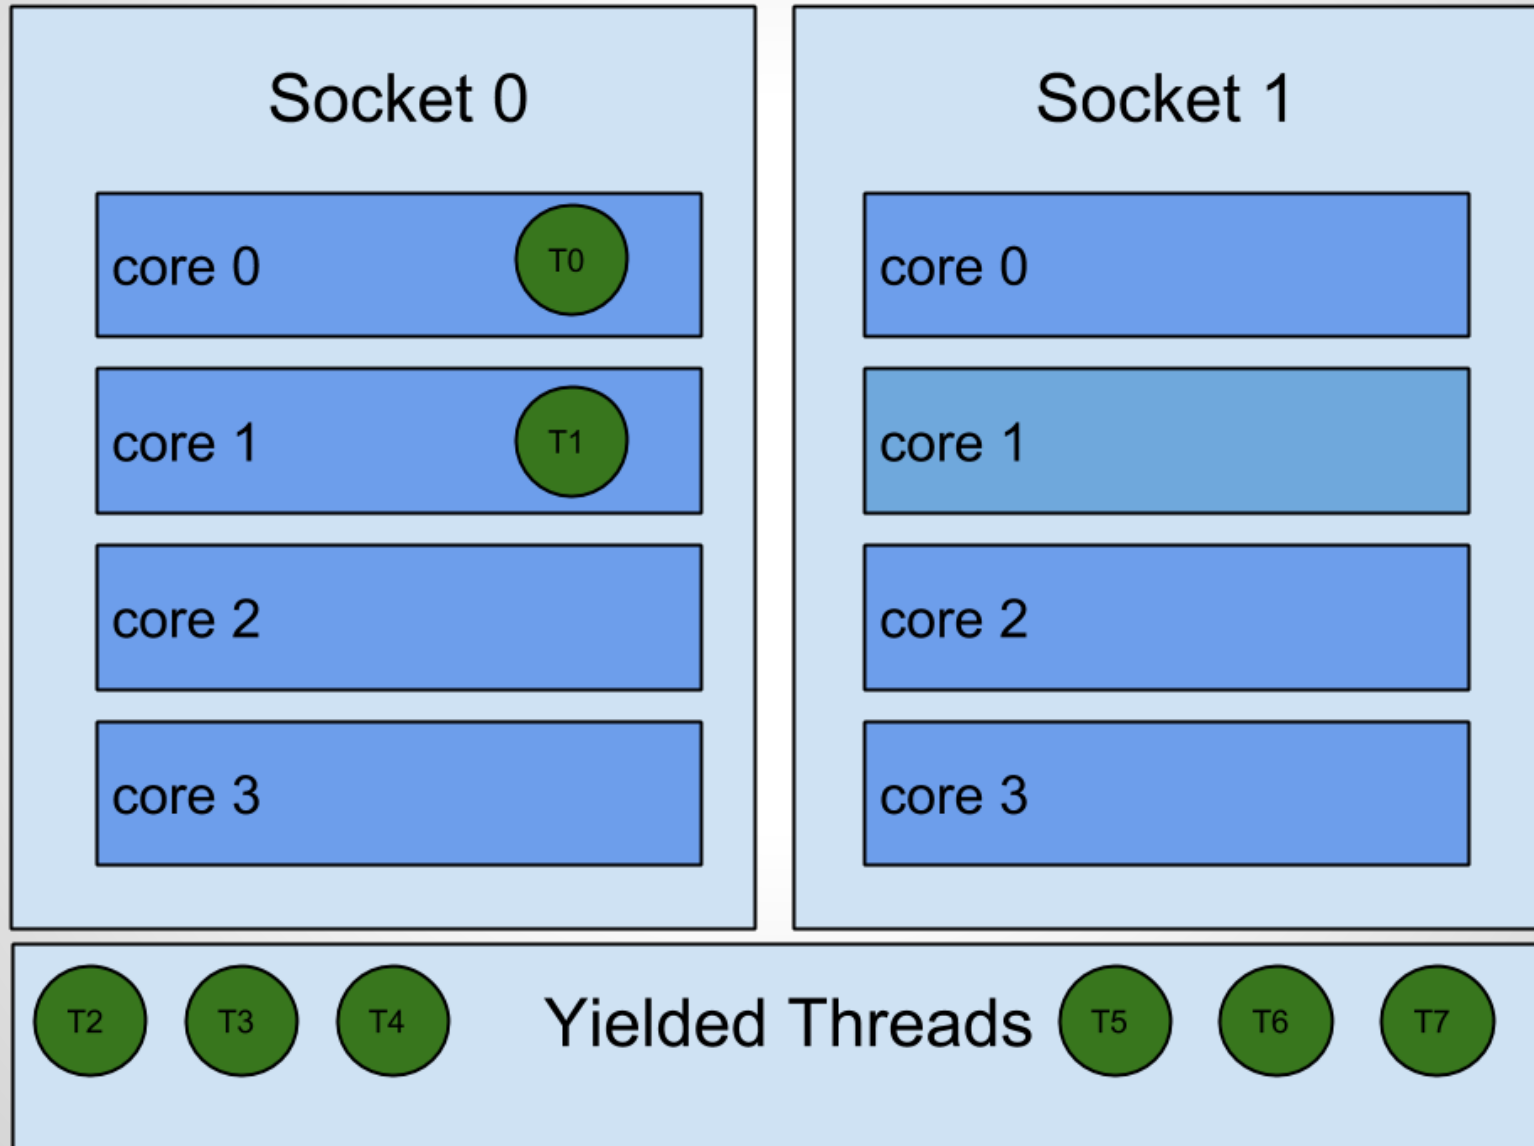
@Override
public void run() {
    while (true) {
        PacketList packets = rp_ind.getBurst();
        if (packets != null) {
            for (Packet p : packets) {
                inspect(p);
            }
        }
    }
}

```

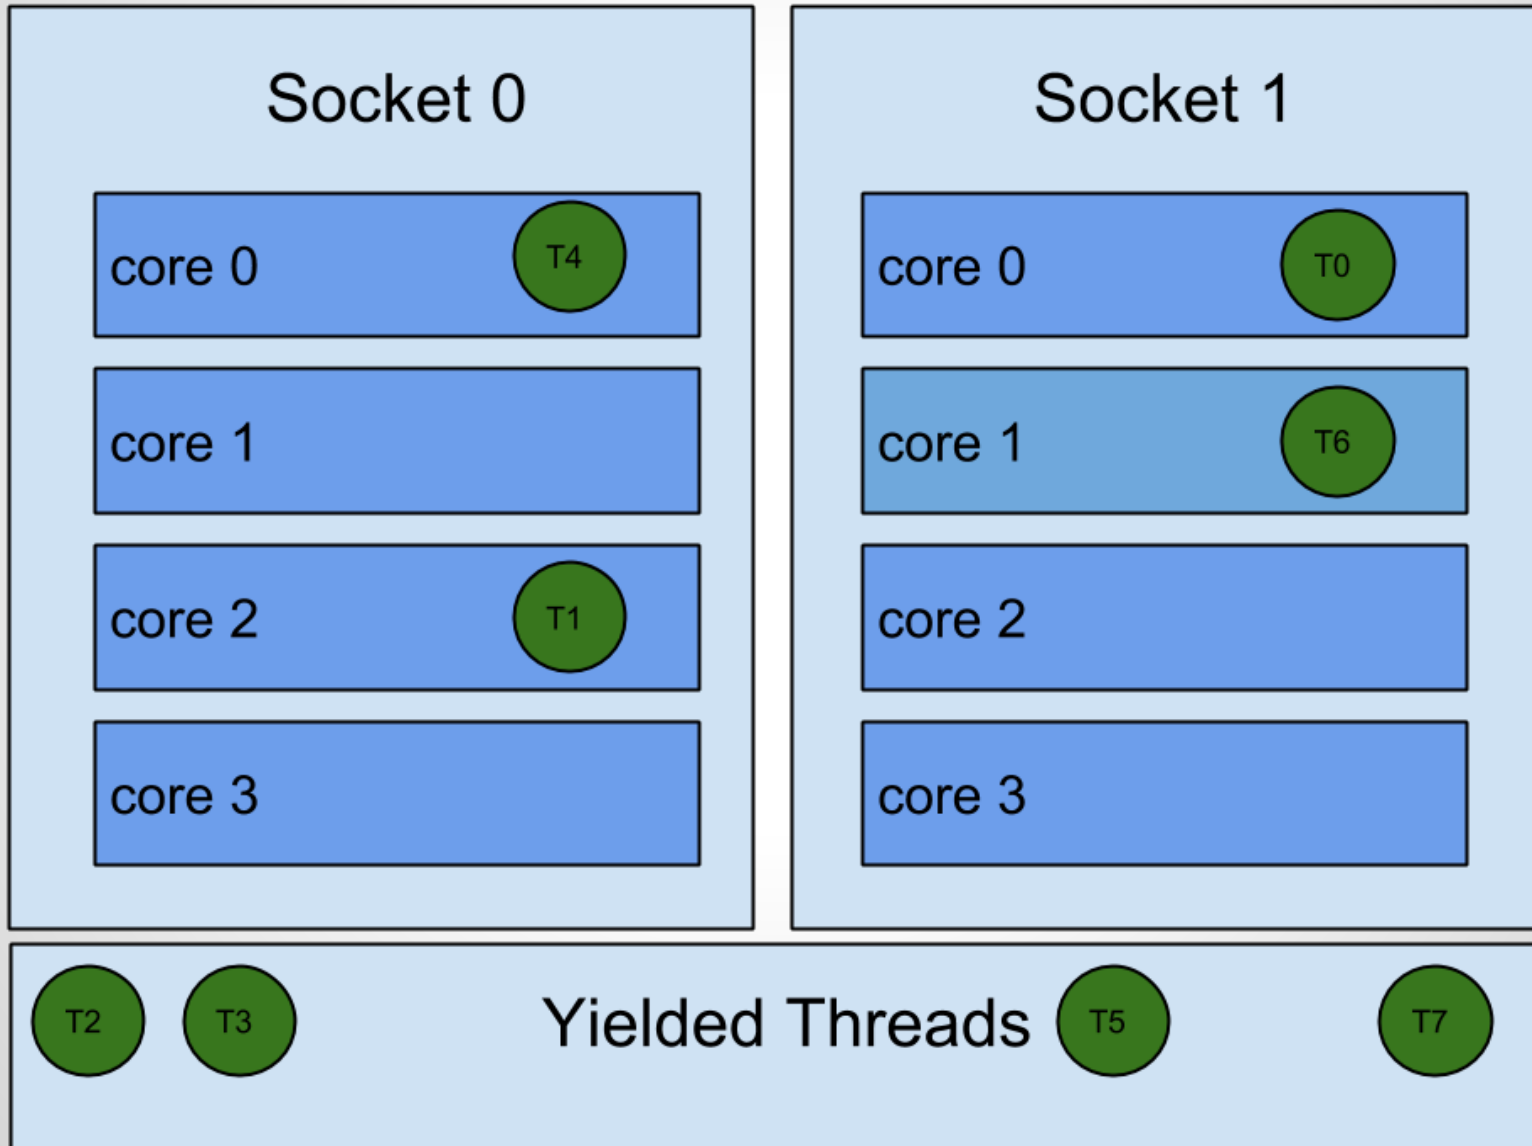
Processor Affinity (CPU Pinning)

- Management of OS thread scheduling
- Stops context switching of threads between cores
- Ideal for continuous looping threads
- For application with little intra-thread communication

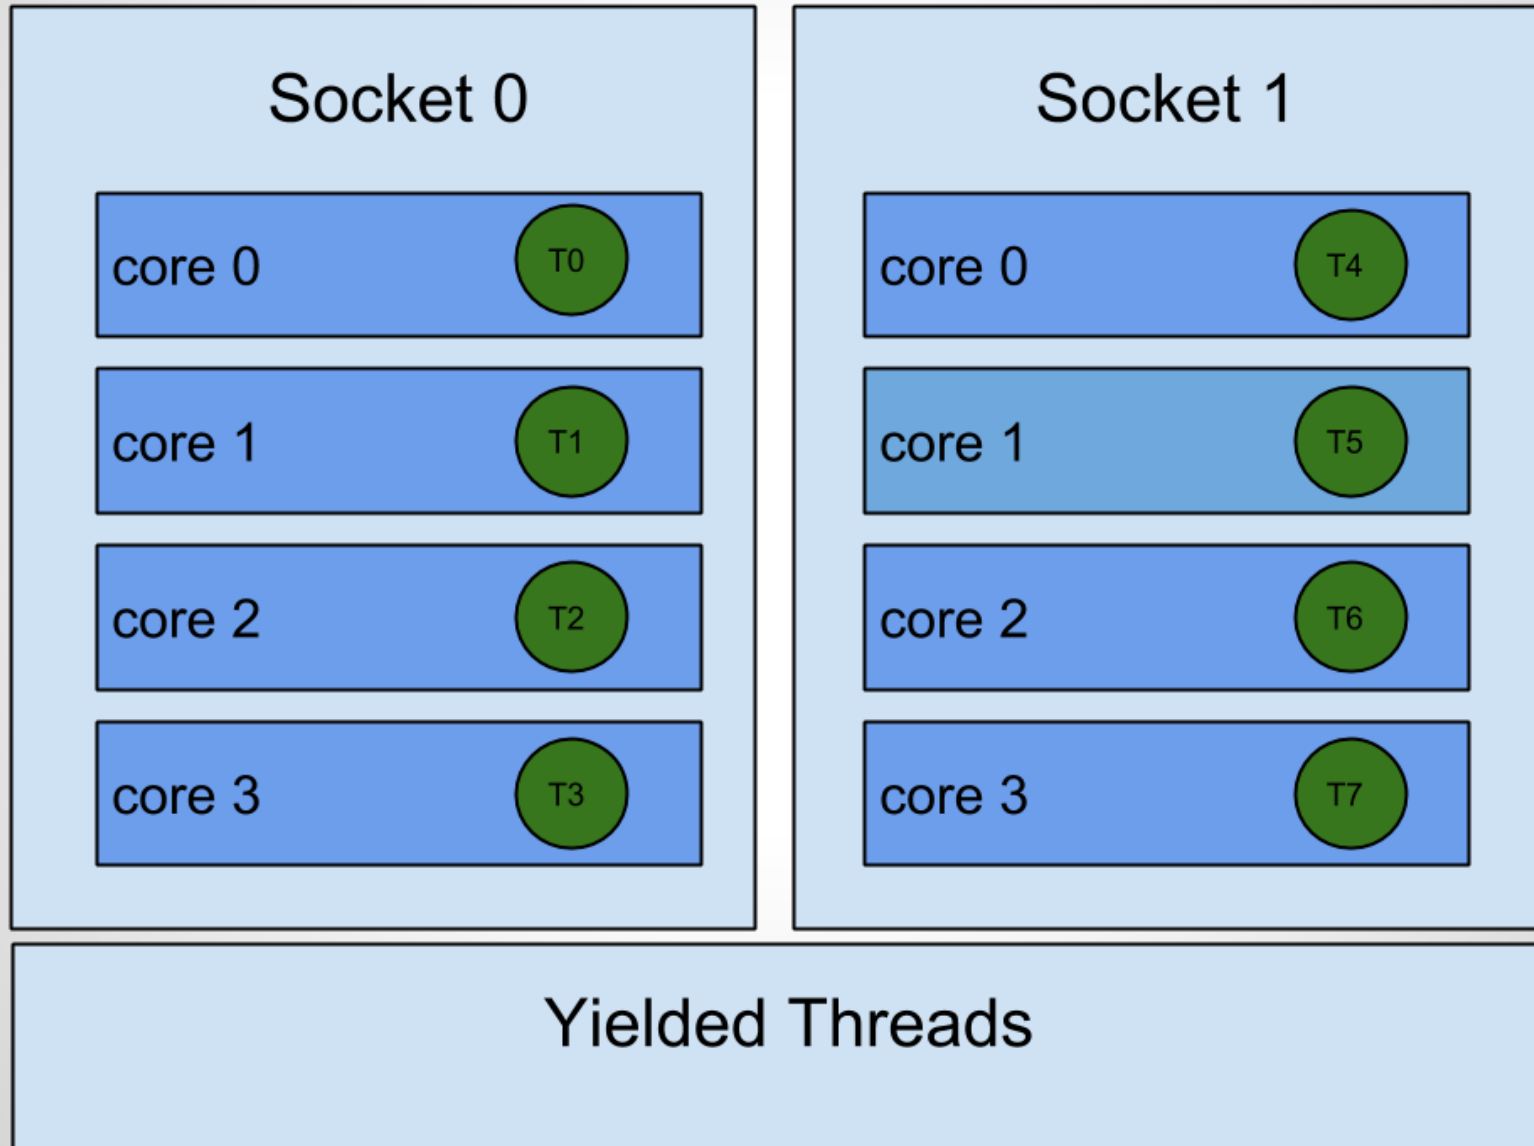
Processor Affinity (CPU Pinning)



Processor Affinity (CPU Pinning)



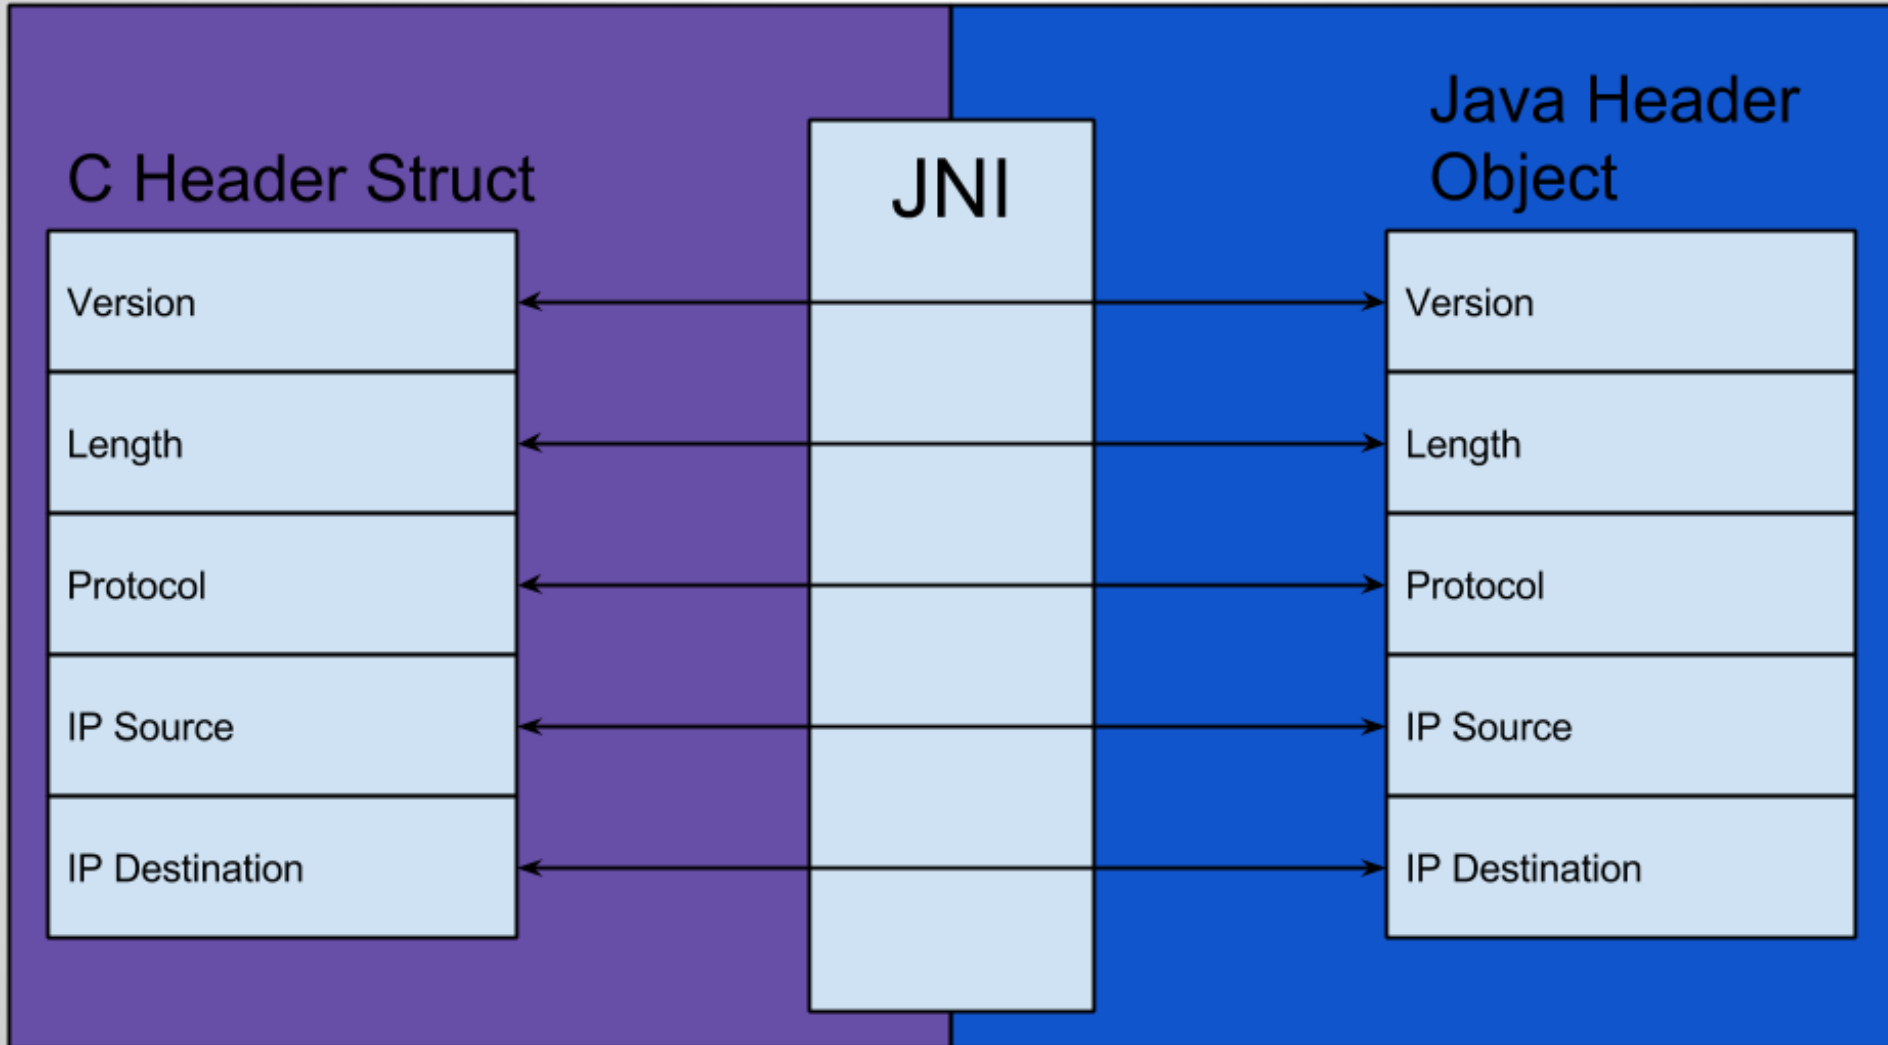
Processor Affinity (CPU Pinning)



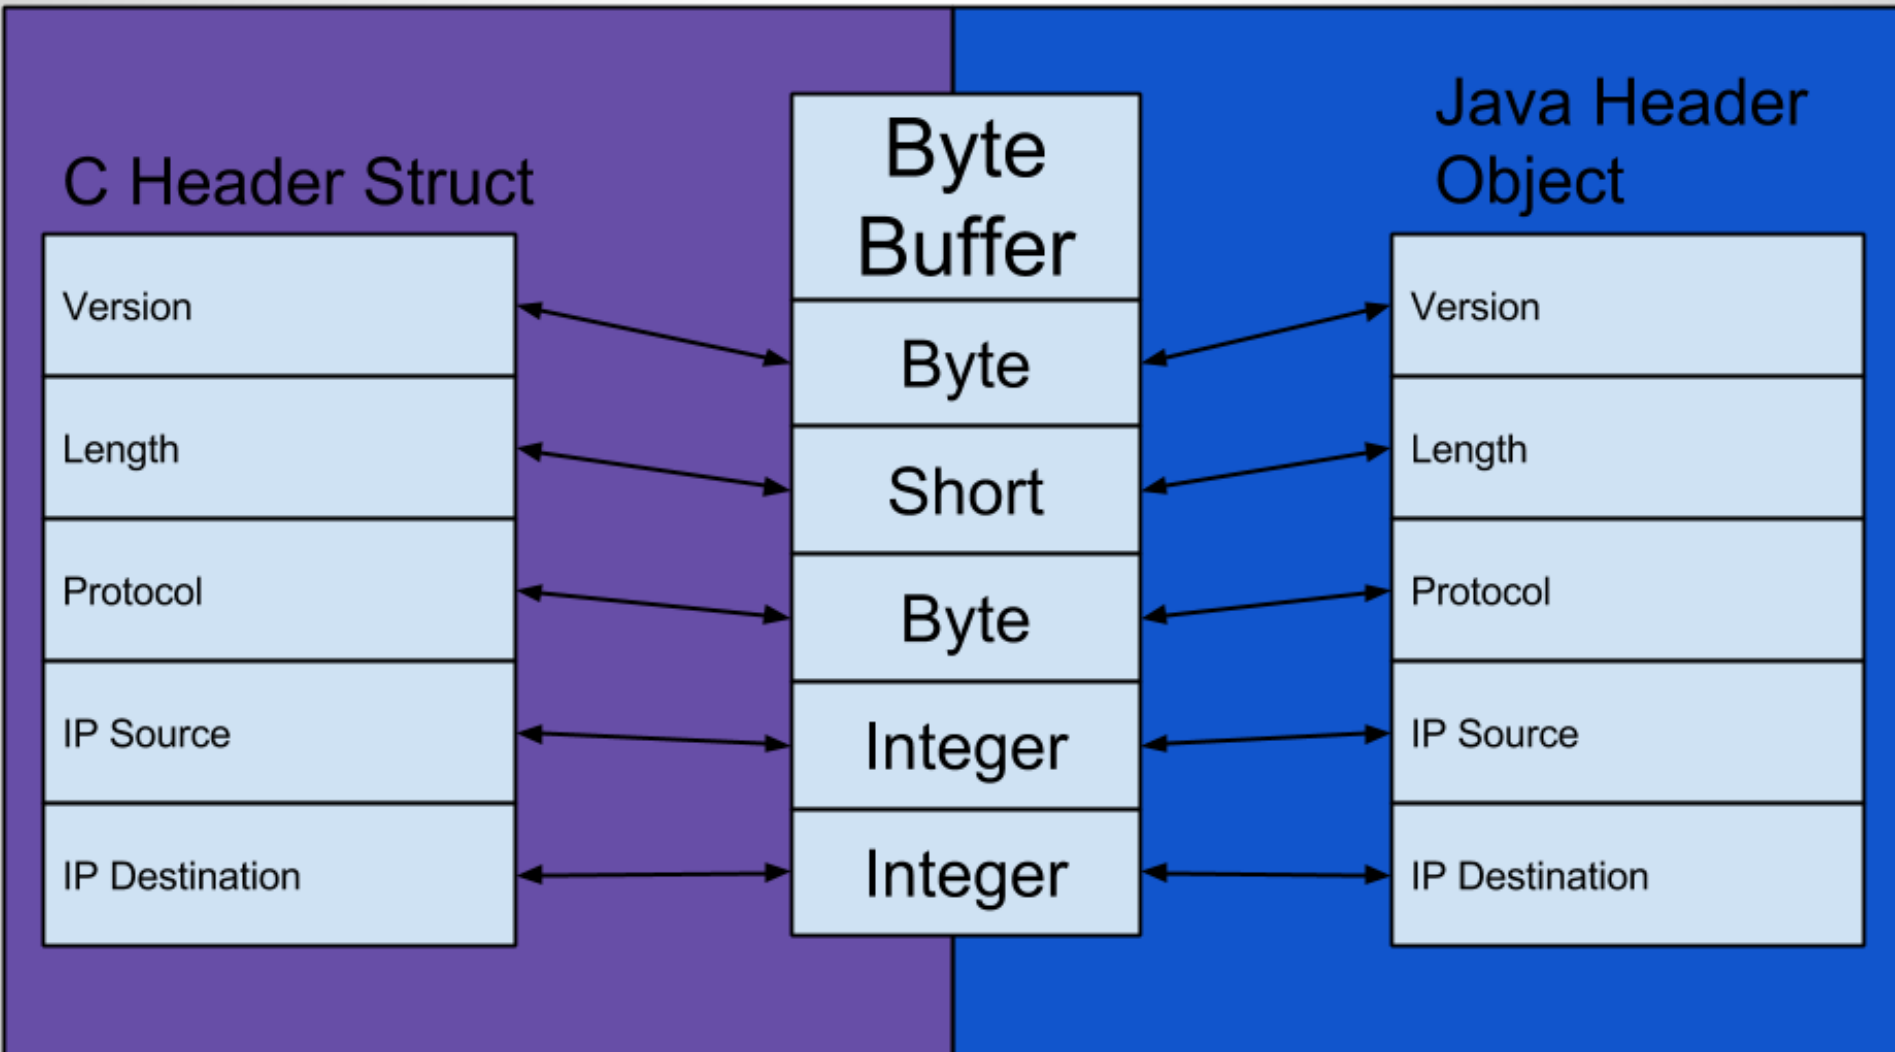
Data Sharing

- Between native and Java code
- Primarily packet meta-data
- Accessing C structs in Java with OO
- 4 possible methods

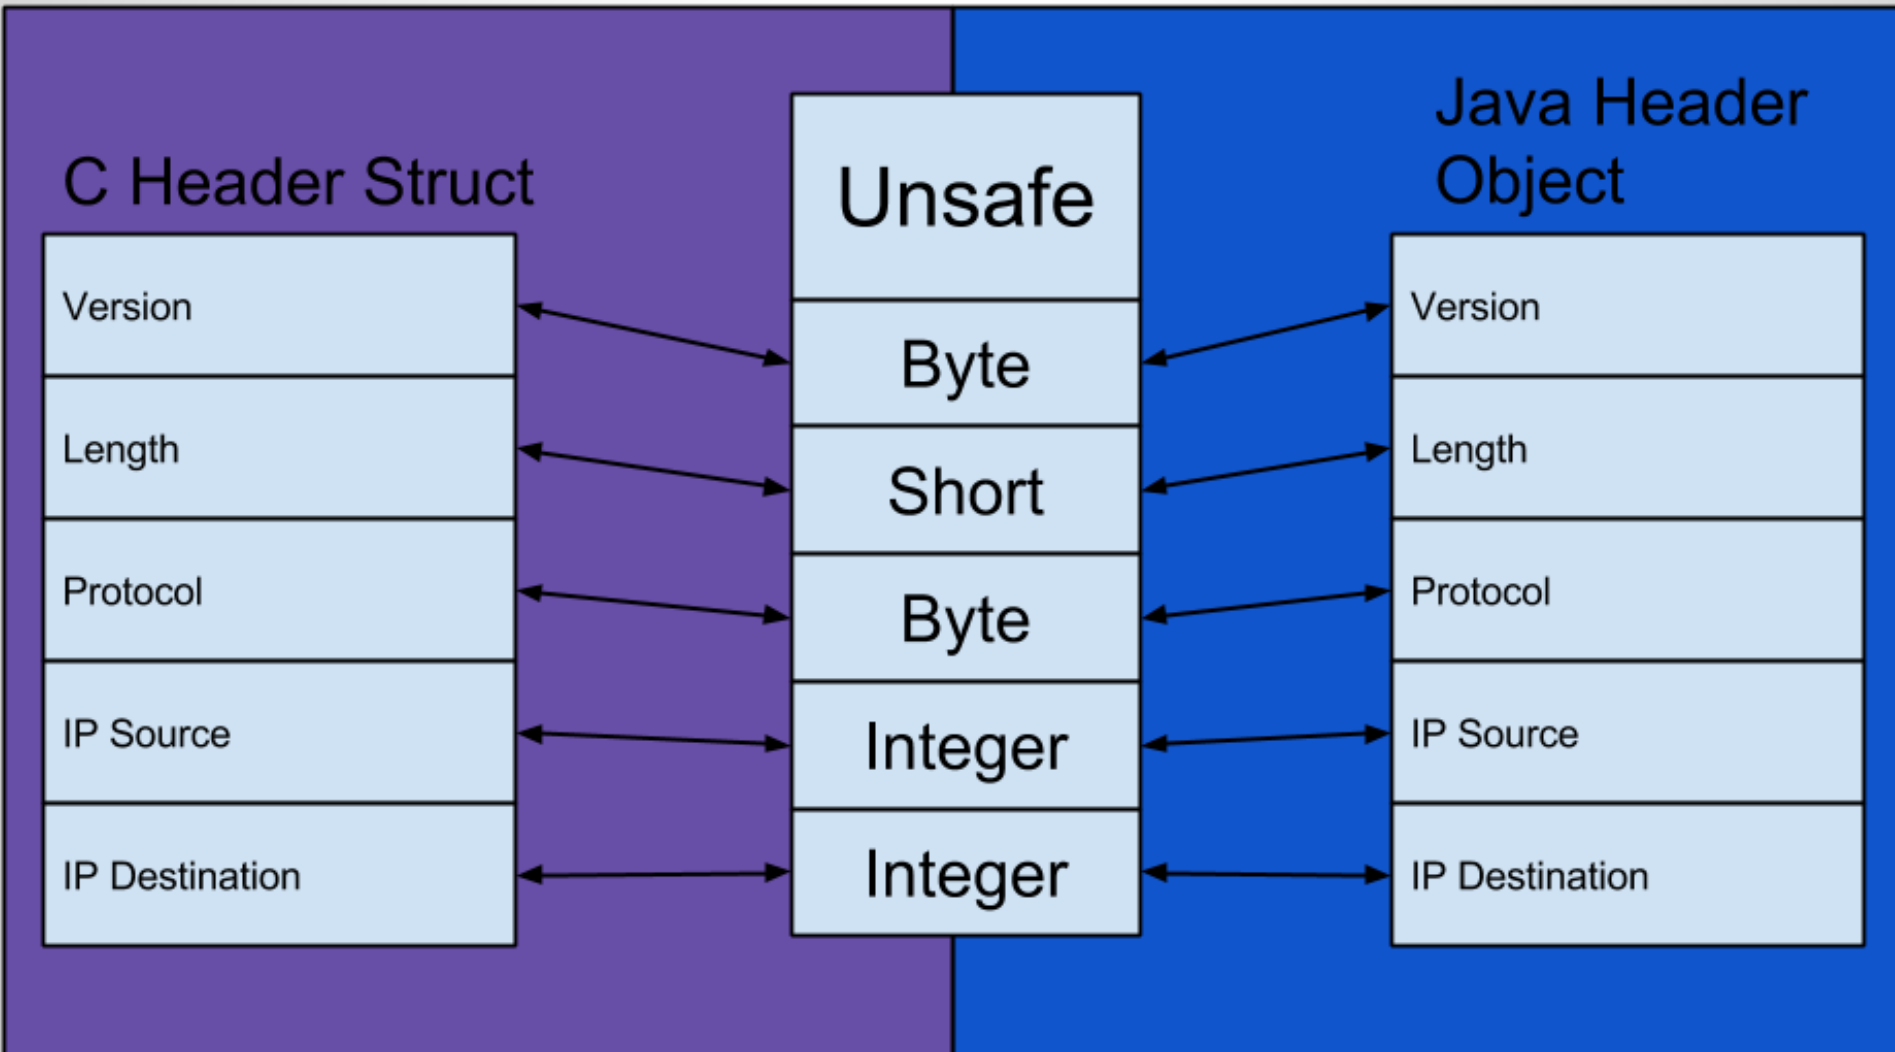
Data Sharing - Object Passing



Data Sharing - Byte Buffers



Data Sharing - Java Unsafe



Data Sharing - Direct Access

C Header Struct

Version

Length

Protocol

IP Source

IP Destination

Unsafe


*mbuf[0]

*ip_hdr[0]


Java Header Object

mbuf_pointer


ip_hdr_pointer



nanoseconds (logarithmic scale)



Implementation - Data Sharing



Implementation - Packet

```
public class Packet {  
  
    long mbuf_pointer;  
    long hdr_pointer;  
    UnsafeAccess ua;  
  
    public long getSrcAddr() {  
        ua.setCurrentPointer(packet_pointer + SRC_ADDR_OFFSET);  
        return ua.getInt();  
    }  
  
    public void setSrcAddr(long src_addr) {  
        ua.setCurrentPointer(packet_pointer + SRC_ADDR_OFFSET);  
        ua.putInt(src_addr);  
    }  
}
```

Unsafe Abstraction

- Handles pointer arithmetic
- Conversions between unsigned and signed
- Bound checking for negatives and overflows
- Endianness

C Struct Packing

```
struct foo {  
    char *p;           // 8 bytes  
    char c;            // 1 byte  
    char pad[7];       // 7 bytes  
    long x;            // 8 bytes  
}
```

C Struct Packing

```
struct __attribute__((packed)) foo {  
    char *p;           // 8 bytes  
    char c;            // 1 byte  
    char pad[7];      // 7 bytes  
    long x;            // 8 bytes  
}
```


C Struct Packing

Advantages

- Reduced Memory
- Allows for packet member pointer arithmetic
- Data usable over different architectures


Disadvantages

- Reduced access speed to members




Performance Testing


- Initially on packet capture to test max speeds
- Use Pktgen software
- Compare C & Java implementation
- Further testing on IP firewall

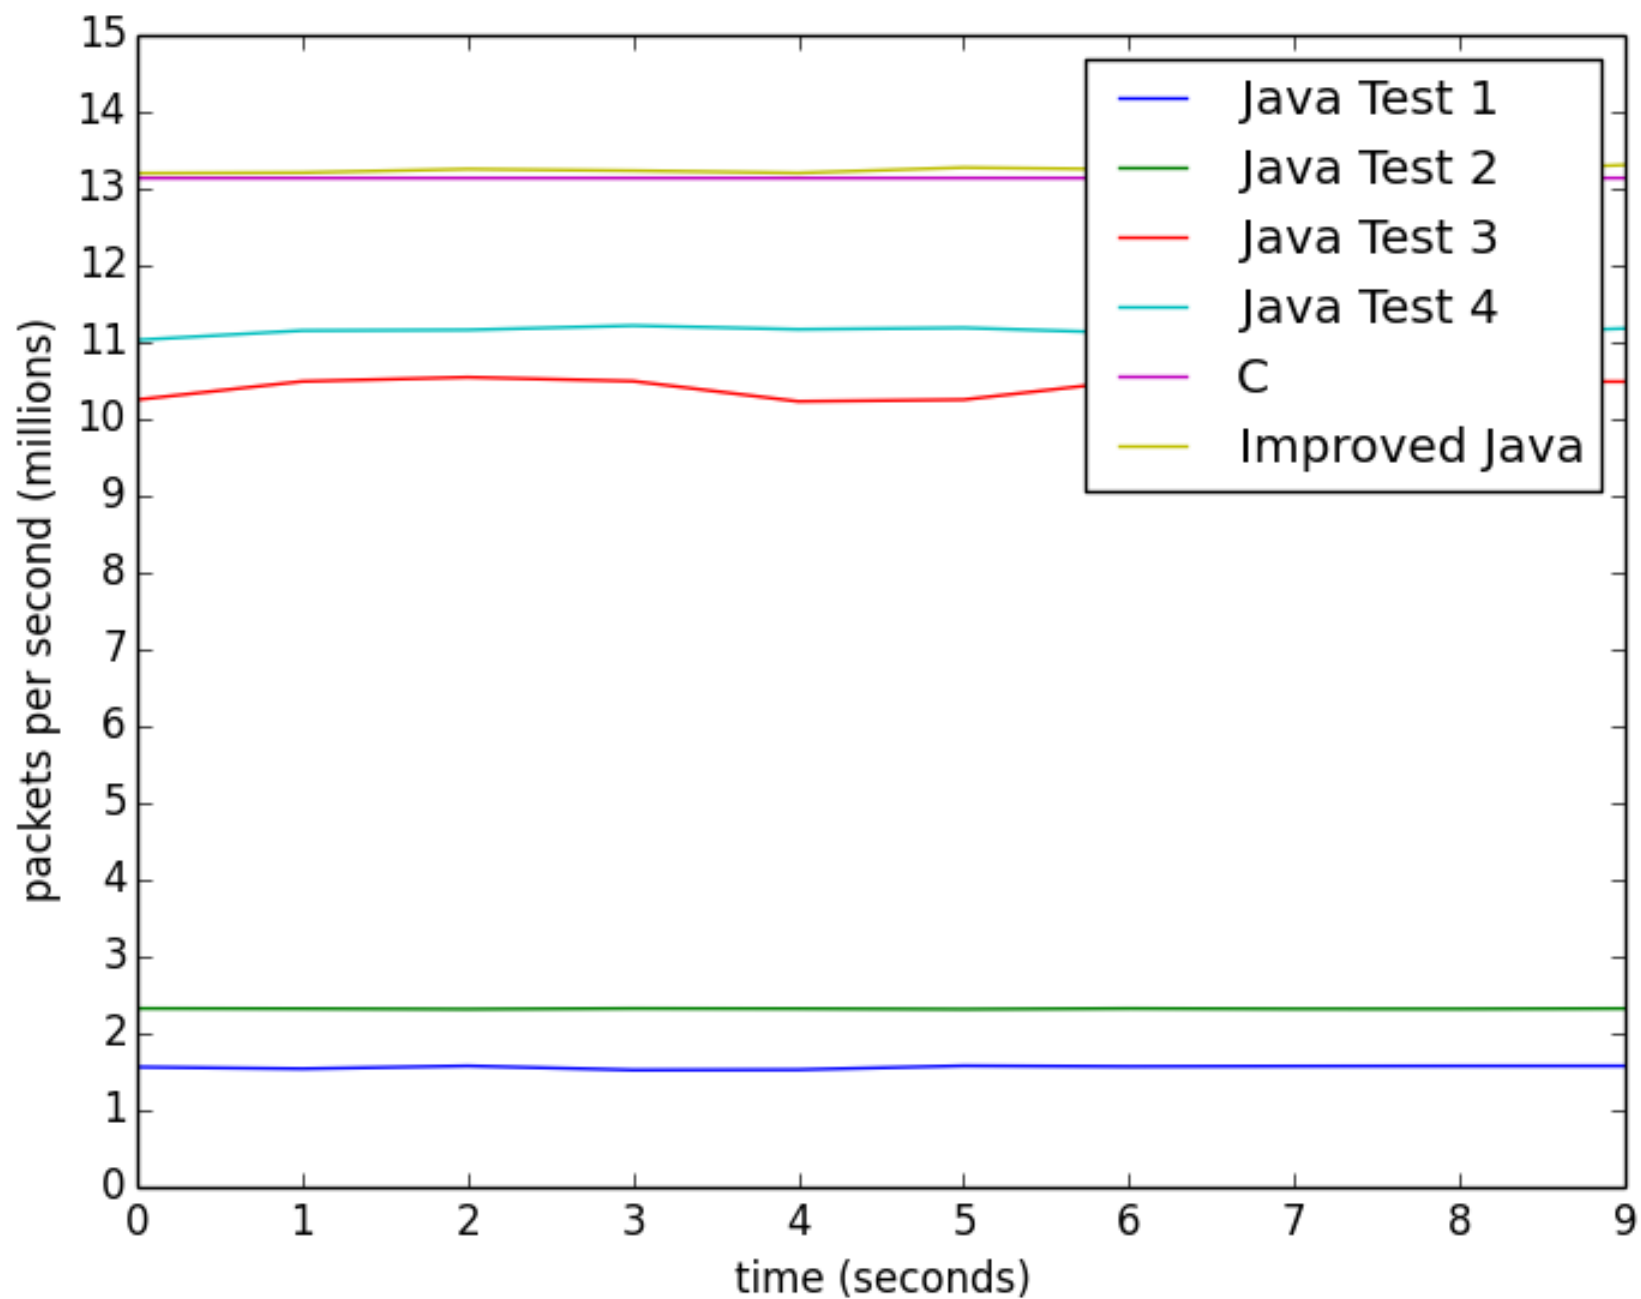



Name	Instance count ▾	Size
java.lang.reflect.Field	457,884	32,967 kB
UnsafeAccess	457,712	14,646 kB
Ipv4Packet	455,620	14,579 kB
java.lang.Object[]	57,772	5,621 kB
java.util.ArrayList	14,424	346 kB
java.util.ArrayList\$SubList	14,237	569 kB
java.util.ArrayList\$Itr	14,237	455 kB
java.lang.String	5,749	137 kB
Ipv6Packet	2,089	66,848 bytes
java.lang.Class	1,053	695 kB


	Hot spot	Inherent time ▾	Average Time	Invocations
⚙️ ⚠️	java.lang.Class.getDeclaredField	3,718 ms (14 %)	2 μs	1,446,197
⚙️ ⚠️	ReceivePoller.getBurst	3,491 ms (13 %)	77 μs	44,986
⚙️ ⚠️	Packet.<init>	2,837 ms (11 %)	1 μs	1,446,197
⚙️ ⚠️	UnsafeAccess.<init>	2,819 ms (11 %)	1 μs	1,446,197
⚙️ ⚠️	Ipv4Packet.getLength	1,197 ms (4 %)	0 μs	2,865,814
⚙️ ⚠️	UnsafeAccess.getShort	1,137 ms (4 %)	0 μs	2,924,090
⚙️ ⚠️	UnsafeAccess.setCurrentPointer	824 ms (3 %)	0 μs	7,294,378
⚙️ ⚠️	PacketFreeer.freePacket	792 ms (3 %)	0 μs	1,439,564
⚙️ ⚠️	UnsafeAccess.getLong	757 ms (2 %)	0 μs	2,879,104
⚙️ ⚠️	PacketFreeer.freeBurst	679 ms (2 %)	15 μs	44,987
⚙️ ⚠️	CaptureProcessor.inspect	591 ms (2 %)	0 μs	1,439,564
⚙️ ⚠️	Ipv4Packet.getVersionIhl	587 ms (2 %)	0 μs	1,439,552
⚙️ ⚠️	UnsafeAccess.getBytes	567 ms (2 %)	0 μs	1,439,552
⚙️ ⚠️	Ipv4Packet.getVersion	440 ms (1 %)	0 μs	1,439,552
⚙️ ⚠️	java.util.List.get	416 ms (1 %)	0 μs	2,879,148
⚙️ ⚠️	java.util.List.add	375 ms (1 %)	0 μs	2,879,115
⚙️ ⚠️	UnsafeAccess.putLong	358 ms (1 %)	0 μs	1,439,584
⚙️ ⚠️	Ipv4Packet.<init>	352 ms (1 %)	0 μs	1,439,552
⚙️ ⚠️	PacketFreeer.isTimedOut	350 ms (1 %)	0 μs	1,394,575
⚙️ ⚠️	sun.misc.Unsafe.getShort	319 ms (1 %)	0 μs	2,924,090
⚙️ ⚠️	sun.misc.Unsafe.getLong	318 ms (1 %)	0 μs	2,879,104
⚙️ ⚠️	Utils.signToUnsign(short)	281 ms (1 %)	0 μs	2,924,090
⚙️ ⚠️	sun.misc.Unsafe.getBytes	275 ms (1 %)	0 μs	1,439,552
⚙️ ⚠️	java.lang.reflect.Field.setAccessible	201 ms (0 %)	0 μs	1,446,197
⚙️ ⚠️	java.lang.System.currentTimeMillis	195 ms (0 %)	0 μs	1,439,562
⚙️ ⚠️	UnsafeAccess.longSize	183 ms (0 %)	0 μs	1,484,538











What's Next?

- Test with pipeline model
- Multiple cores
- Full-scale application in system
- Add further DPDK features to DPDK-Java

Thanks For Listening!

Any Questions?